

```
In [1]: from custom_class import CustomCrossValidator
        from tools import *

        import pandas as pd
        import numpy as np
        import seaborn as sns
        import matplotlib.pyplot as plt

        from sklearn.linear_model import LogisticRegression, Lasso
        from sklearn.model_selection import train_test_split
        from sklearn.feature_selection import RFE, VarianceThreshold
        from sklearn.preprocessing import StandardScaler, OneHotEncoder, LabelEncoder
        from sklearn.model_selection import cross_val_score, StratifiedKFold
        from sklearn.tree import DecisionTreeClassifier
        from statsmodels.stats.outliers_influence import variance_inflation_factor
        from sklearn.metrics import recall_score, confusion_matrix

        import folium
        from folium.plugins import MarkerCluster
```

```
In [2]: # Import the split data set
        target_df = pd.read_csv("data/target_columns.csv")
        independent_df = pd.read_csv("data/independent_columns.csv")
```

```
In [3]: # Joining the data sets
        wells_df = pd.merge(target_df, independent_df, on='id', how='inner')
```

```
In [4]: # Standardize categorical variables
        wells_df = wells_df.applymap(lambda x: x.lower() if isinstance(x, str) else x)
```

```
In [5]: # Convert boolean values to strings
        wells_df = wells_df.applymap(lambda x: str(x) if isinstance(x, bool) else x)
```

```
In [6]: # Standardize missing values to NaN
        missing_values = ['', ' ', 'na', 'n/a', 'unknown', 'other', 'none']
        wells_df.replace(missing_values, np.nan, inplace=True)
```

```
In [7]: # Set 'id' as the index of the DataFrame
        wells_df.set_index('id', inplace=True)
```

```
In [8]: wells_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 59400 entries, 69572 to 26348
Data columns (total 40 columns):
#   Column                Non-Null Count  Dtype
---  -
0   status_group          59400 non-null  object
1   amount_tsh            59400 non-null  float64
2   date_recorded         59400 non-null  object
3   funder                55759 non-null  object
4   gps_height            59400 non-null  int64
5   installer             55741 non-null  object
6   longitude             59400 non-null  float64
7   latitude              59400 non-null  float64
8   wpt_name              55832 non-null  object
9   num_private           59400 non-null  int64
10  basin                 59400 non-null  object
11  subvillage            59029 non-null  object
12  region                59400 non-null  object
13  region_code           59400 non-null  int64
14  district_code         59400 non-null  int64
15  lga                   59400 non-null  object
16  ward                  59400 non-null  object
17  population            59400 non-null  int64
18  public_meeting        56066 non-null  object
19  recorded_by           59400 non-null  object
20  scheme_management     54756 non-null  object
21  scheme_name           30565 non-null  object
22  permit               56344 non-null  object
23  construction_year     59400 non-null  int64
24  extraction_type       52970 non-null  object
25  extraction_type_group  52970 non-null  object
26  extraction_type_class  52970 non-null  object
27  management            57995 non-null  object
28  management_group      57896 non-null  object
29  payment               50189 non-null  object
30  payment_type          50189 non-null  object
31  water_quality         57524 non-null  object
32  quality_group         57524 non-null  object
33  quantity              58611 non-null  object
34  quantity_group        58611 non-null  object
35  source                59122 non-null  object
36  source_type           59122 non-null  object
37  source_class          59122 non-null  object
38  waterpoint_type       53020 non-null  object
39  waterpoint_type_group 53020 non-null  object
dtypes: float64(3), int64(6), object(31)
memory usage: 18.6+ MB
```

## Target Variable Investigation

```
In [9]: wells_df['status_group'].value_counts(normalize=True)
```

```
Out[9]: functional                0.543081
non functional                  0.384242
functional needs repair         0.072677
Name: status_group, dtype: float64
```

There is a class imbalance between the three classes inside of the target variable. We will need to handle this through by either oversampling or SMOTE'ing to bring the minority class in proportion with the majority class.

Given the small proportion of the 'function needs repair' class we will combine the 'functional needs repair' and 'non functional' class together to create the new class 'needs attention'. Then we will cast the 'functional' class of wells as 'does not need attention'.

```
In [10]: # Define the mapping of the new labels
label_mapping = {
    'functional': 0,
    'non functional': 1,
    'functional needs repair': 1
}

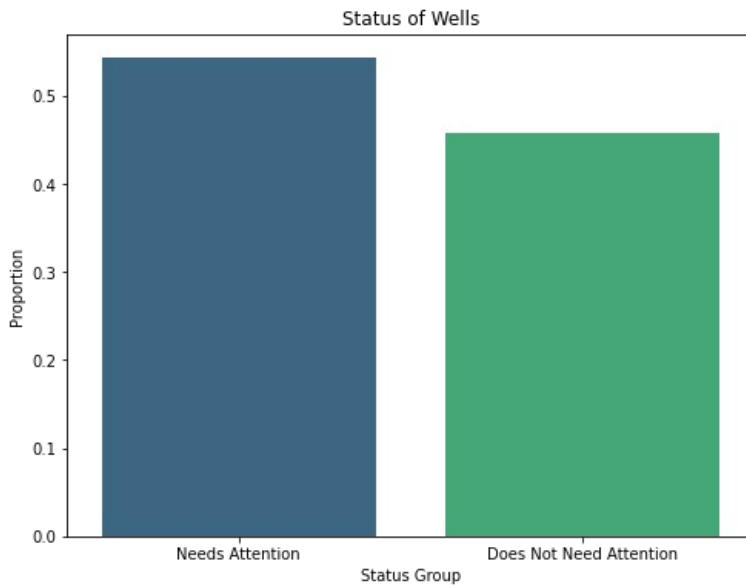
# Replace the values in the 'status_group' column
wells_df['status_group'] = wells_df['status_group'].replace(label_mapping)
```

```
In [11]: status_counts_normalized = wells_df['status_group'].value_counts(normalize=True)
status_counts_normalized
```

```
Out[11]: 0    0.543081
1    0.456919
Name: status_group, dtype: float64
```

```
In [12]: # Plotting
```

```
plt.figure(figsize=(8, 6))
sns.barplot(x=status_counts_normalized.index, y=status_counts_normalized.values, palette='viridis')
plt.xlabel('Status Group')
plt.ylabel('Proportion')
plt.title('Status of Wells')
plt.xticks([0, 1], ['Needs Attention', 'Does Not Need Attention'])
plt.show()
```



## Data Set Visualizations

### Map Visualization for Wells Status

```
In [13]: # Create a map centered around the mean latitude and longitude
map_center = [wells_df['latitude'].mean(), wells_df['longitude'].mean()]
mymap = folium.Map(location=map_center, zoom_start=6)
```

```
In [14]: # Define a function to determine marker color based on status_group
def get_color(status):
    return 'green' if status == 0 else 'red'
```

```
In [15]: # Create markers for each row in the DataFrame
for index, row in wells_df.iterrows():
    color = get_color(row['status_group'])

    folium.CircleMarker(
        location=[row['latitude'], row['longitude']],
        radius=5,
        color=color,
        fill=True,
        fill_color=color
    ).add_to(mymap)
```

```
In [16]: # Display the map
mymap
```

Out [16]



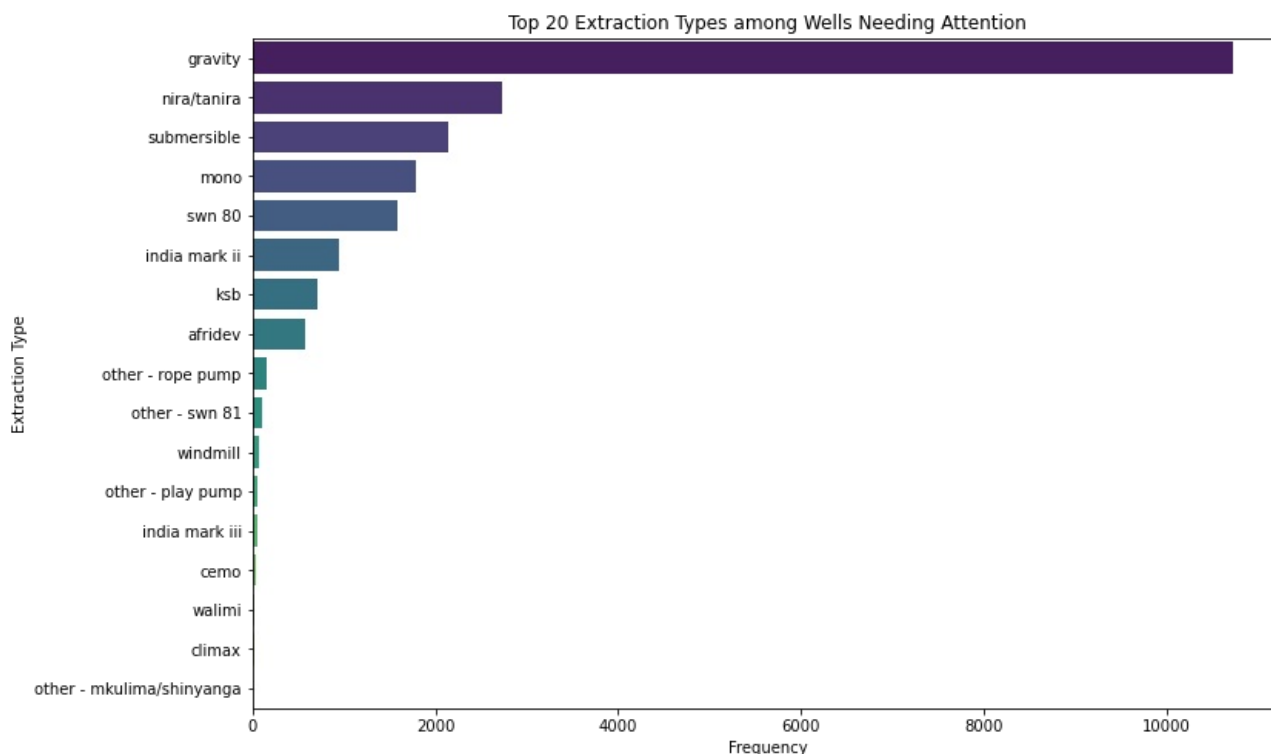
ap: File -> Trust Notebook

## Extraction Type Investigation for Wells Needing Attention

```
In [17]: # Counting the frequency of each extraction_type for wells needing attention (status_group == 1)
attention_wells = wells_df[wells_df['status_group'] == 1]
extraction_type_counts = attention_wells['extraction_type'].value_counts()

# Selecting top 20 categories
top_20_extraction_types = extraction_type_counts.head(20)
```

```
In [18]: # Plotting
plt.figure(figsize=(12, 8))
sns.barplot(x=top_20_extraction_types.values, y=top_20_extraction_types.index, palette='viridis')
plt.xlabel('Frequency')
plt.ylabel('Extraction Type')
plt.title('Top 20 Extraction Types among Wells Needing Attention')
plt.show()
```

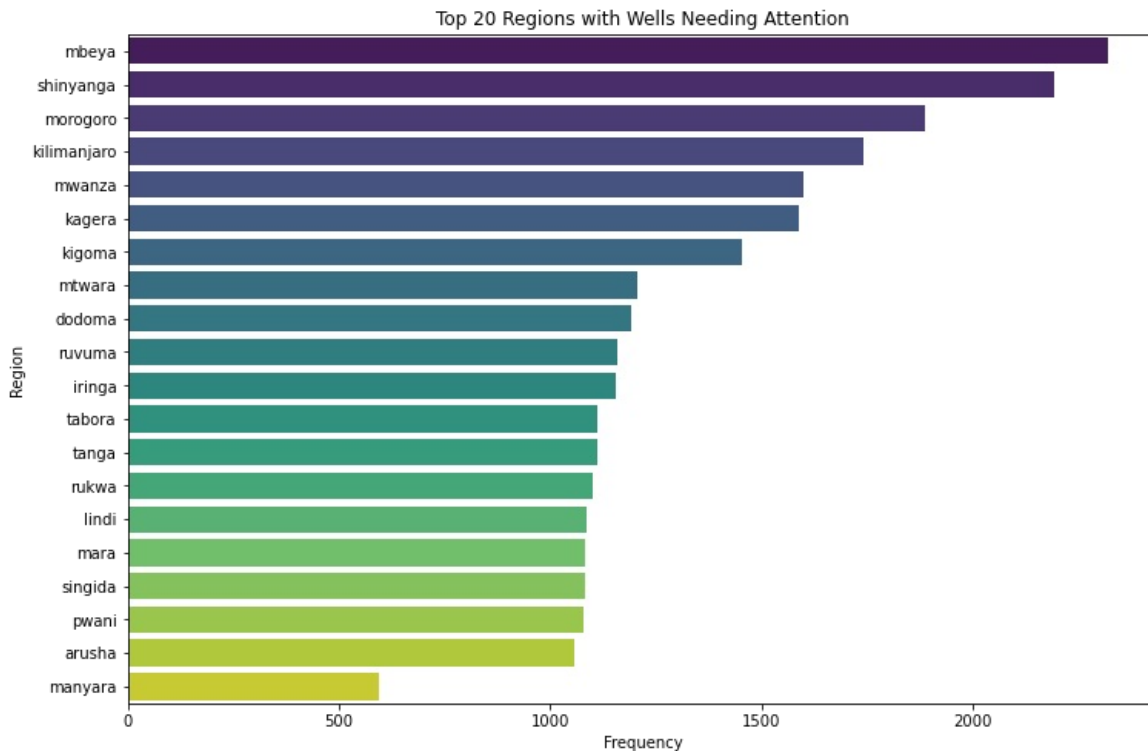


## Region Invetigation for Wells Needing Attention

```
In [19]: # Counting the frequency of each region for wells needing attention (status_group == 1)
region_counts = attention_wells['region'].value_counts()

# Selecting top 20 regions
top_20_regions = region_counts.head(20)
```

```
In [20]: # Plotting
plt.figure(figsize=(12, 8))
sns.barplot(x=top_20_regions.values, y=top_20_regions.index, palette='viridis')
plt.xlabel('Frequency')
plt.ylabel('Region')
plt.title('Top 20 Regions with Wells Needing Attention')
plt.show()
```



## Data Cleaning

### Amount TSH: Illogical Values

```
In [21]: filtered_df = wells_df[wells_df['amount_tsh'] >= 50000]

filtered_df['amount_tsh'].value_counts()
```

```
Out[21]: 117000.0    7
50000.0    4
100000.0    3
250000.0    1
138000.0    1
70000.0    1
350000.0    1
60000.0    1
170000.0    1
200000.0    1
120000.0    1
Name: amount_tsh, dtype: int64
```

Most total static head measurements are in the range of hundreds to tens of thousands.

Depths exceeding 50,000 units (e.g., feet or meters) are rare or unusual so we will filter these out as they are outliers.

Most residential wells range from a few meters to around 100 meters (10 to 300 feet).

Wells used for industrial or agricultural purposes may have deeper total static heads, often ranging from tens to several hundred meters (hundreds to thousands of feet).

```
In [22]: wells_df = wells_df[wells_df['amount_tsh'] <= 50000]
```

### Longitude and Latitude: Zero Value Coordinates

```
In [23]: wells_df['longitude'].value_counts()
```

```
Out[23]: 0.000000    1812
         37.252194      2
         37.297680      2
         33.010510      2
         39.093484      2
         ...
         36.871976      1
         37.579803      1
         33.196490      1
         34.017119      1
         30.163579      1
         Name: longitude, Length: 57498, dtype: int64
```

```
In [24]: wells_df['latitude'].value_counts()
```

```
Out[24]: -2.000000e-08    1812
         -7.104923e+00      2
         -6.964258e+00      2
         -6.963565e+00      2
         -7.056372e+00      2
         ...
         -3.411358e+00      1
         -8.958207e+00      1
         -3.261113e+00      1
         -5.435762e+00      1
         -2.598965e+00      1
         Name: latitude, Length: 57499, dtype: int64
```

```
In [25]: # Filtering to see if any other rows are affected
         filtered_df = wells_df[wells_df['longitude'] == 0]

         print(len(filtered_df))

         filtered_df.head()
```

1812

	status_group	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude	wpt_name	num_priv
id										
6091	0	0.0	2013-02-10	dwsp	0	dwe	0.0	-2.000000e-08	muungano	
32376	1	0.0	2011-08-01	government of tanzania	0	government	0.0	-2.000000e-08	polisi	
72678	0	0.0	2013-01-30	wvt	0	wvt	0.0	-2.000000e-08	wvt tanzania	
56725	1	0.0	2013-01-17	netherlands	0	dwe	0.0	-2.000000e-08	kikundi cha wakina mama	
13042	1	0.0	2012-10-29	hesawa	0	dwe	0.0	-2.000000e-08	kwakisusi	

5 rows × 40 columns



Since these wells have longitude and latitude of 0 it is safe to say that these wells while they were recorded either have an unknown location or do not have a location inside of Tanzania and as such are not reliable data point.

We will be dropping all of these rows from the dataset.

```
In [26]: wells_df = wells_df[wells_df['longitude'] != 0]
```

## Installer Column: Typos

```
In [27]: wells_df['installer'].nunique()
```

```
Out[27]: 1903
```

```
In [28]: installer_replacement_dict = {
         'central government': 'central government',
         'tanzania government': 'central government',
         'cental government': 'central government',
         'cebral government': 'central government',
         'centra government': 'central government',
         }
```

'central govt': 'central government',  
'centr': 'central government',  
'centra govt': 'central government',  
'tanzanian government': 'central government',  
'tanzania': 'central government',  
  
'district council': 'district council',  
'counc': 'district council',  
'district counci': 'district council',  
'council': 'district council',  
'coun': 'district council',  
'distri': 'district council',  
'district council': 'district council',  
  
'villigers': 'villagers',  
'villager': 'villagers',  
'villa': 'villagers',  
'village': 'villagers',  
'villi': 'villagers',  
'village council': 'villagers',  
'village counil': 'villagers',  
'villages': 'villagers',  
'vill': 'villagers',  
'village community': 'villagers',  
'villaers': 'villagers',  
'villag': 'villagers',  
'villege council': 'villagers',  
'villagerd': 'villagers',  
'village technician': 'villagers',  
'village office': 'villagers',  
'village community members': 'villagers',  
'village government': 'villagers',  
'village govt': 'villagers',  
  
'district water department': 'district water department',  
'district water depar': 'district water department',  
'distric water department': 'district water department',  
  
'finw': 'fini water',  
'fini water': 'fini water',  
'fin water': 'fini water',  
'finwater': 'fini water',  
'finn water': 'fini water',  
'fw': 'fini water',  
  
'rc church': 'rc church',  
'rc churc': 'rc church',  
'rc': 'rc church',  
'rc ch': 'rc church',  
'rc c': 'rc church',  
'rc cathoric': 'rc church',  
'ch': 'rc church',  
  
'world vision': 'world vision',  
'world division': 'world vision',  
'world viission': 'world vision',  
  
'unisef': 'unicef',  
'unicef': 'unicef',  
  
'danid': 'danida',  
  
'commu': 'community',  
'communit': 'community',  
'adra /community': 'community',  
'adra/community': 'community',  
'adra/ community': 'community',  
  
'government': 'government',  
'gover': 'government',  
'governme': 'government',  
'govern': 'government',  
'govern': 'government',  
'gove': 'government',  
'governmen': 'government',  
  
'hesawa': 'hesawa',  
  
'jaica': 'jaica',  
'jica': 'jaica',  
'jeica': 'jaica',  
'jaica co': 'jaica',

```
'kkkt _ konde and dwe': 'kkkt',
'kkt': 'kkkt',
'kkkt church': 'kkkt'
}
```

```
In [29]: wells_df['installer'] = wells_df['installer'].replace(installer_replacement_dict)
```

```
In [30]: wells_df['installer'].nunique()
```

```
Out[30]: 1835
```

## Formatting Cleanup

```
In [31]: # Function to replace "other - " prefix
def replace_other_prefix(value):
    if isinstance(value, str) and value.startswith('other -'):
        return value.split('other -', 1)[1].strip()
    return value

# Apply the function to the entire DataFrame
wells_df = wells_df.applymap(replace_other_prefix)
```

## Missing Values

### Categorical columns with 0's

```
In [32]: # Replace all 0's with NaNs
for col in wells_df.columns:
    if wells_df[col].dtype == '0':
        wells_df[col] = wells_df[col].replace('0', np.nan)
```

## Applying the correct data types

The columns listed below act more as labels than as integers. They identify a specific aspect of about the row entry and as such should be considered categorical variables.

```
In [33]: wells_df['region_code'] = wells_df['region_code'].astype(str)
wells_df['district_code'] = wells_df['district_code'].astype(str)
```

## Missing values

```
In [34]: # Assuming you already have 'expected_length' and 'wells_df' defined
expected_length = 59400
columns_with_missing_values = {}

# Iterate over columns to find those not meeting expected length
for col in wells_df.columns:
    if wells_df[col].count() != expected_length:
        columns_with_missing_values[col] = {
            'missing_count': expected_length - wells_df[col].count(),
            'unique_count': len(wells_df[col].value_counts()),
            'dtype': wells_df[col].dtype # Add data type information
        }

# Print the dictionary
for col, info in columns_with_missing_values.items():
    print(f"{col}': Missing Values: {info['missing_count']}, Unique Values: {info['unique_count']}, Dtype: {info['dtype']}")
```



```
status_group': Missing Values: 1830, Unique Values: 2, Dtype: int64
amount_tsh': Missing Values: 1830, Unique Values: 88, Dtype: float64
date_recorded': Missing Values: 1830, Unique Values: 353, Dtype: object
funder': Missing Values: 6234, Unique Values: 1854, Dtype: object
gps_height': Missing Values: 1830, Unique Values: 2428, Dtype: int64
installer': Missing Values: 6246, Unique Values: 1834, Dtype: object
longitude': Missing Values: 1830, Unique Values: 57497, Dtype: float64
latitude': Missing Values: 1830, Unique Values: 57498, Dtype: float64
wpt_name': Missing Values: 5327, Unique Values: 36707, Dtype: object
num_private': Missing Values: 1830, Unique Values: 65, Dtype: int64
basin': Missing Values: 1830, Unique Values: 9, Dtype: object
subvillage': Missing Values: 2201, Unique Values: 18566, Dtype: object
region': Missing Values: 1830, Unique Values: 21, Dtype: object
region_code': Missing Values: 1830, Unique Values: 27, Dtype: object
district_code': Missing Values: 1830, Unique Values: 20, Dtype: object
lga': Missing Values: 1830, Unique Values: 124, Dtype: object
ward': Missing Values: 1830, Unique Values: 2033, Dtype: object
population': Missing Values: 1830, Unique Values: 1047, Dtype: int64
public_meeting': Missing Values: 4806, Unique Values: 2, Dtype: object
recorded_by': Missing Values: 1830, Unique Values: 1, Dtype: object
scheme_management': Missing Values: 6346, Unique Values: 10, Dtype: object
scheme_name': Missing Values: 29189, Unique Values: 2536, Dtype: object
permit': Missing Values: 4884, Unique Values: 2, Dtype: object
construction_year': Missing Values: 1830, Unique Values: 55, Dtype: int64
extraction_type': Missing Values: 7990, Unique Values: 17, Dtype: object
extraction_type_group': Missing Values: 7990, Unique Values: 12, Dtype: object
extraction_type_class': Missing Values: 7990, Unique Values: 6, Dtype: object
management': Missing Values: 3221, Unique Values: 10, Dtype: object
management_group': Missing Values: 3320, Unique Values: 3, Dtype: object
payment': Missing Values: 10386, Unique Values: 5, Dtype: object
payment_type': Missing Values: 10386, Unique Values: 5, Dtype: object
water_quality': Missing Values: 3491, Unique Values: 7, Dtype: object
quality_group': Missing Values: 3491, Unique Values: 5, Dtype: object
quantity': Missing Values: 2603, Unique Values: 4, Dtype: object
quantity_group': Missing Values: 2603, Unique Values: 4, Dtype: object
source': Missing Values: 2096, Unique Values: 8, Dtype: object
source_type': Missing Values: 2096, Unique Values: 6, Dtype: object
source_class': Missing Values: 2096, Unique Values: 2, Dtype: object
waterpoint_type': Missing Values: 7997, Unique Values: 6, Dtype: object
waterpoint_type_group': Missing Values: 7997, Unique Values: 5, Dtype: object
```

## Filling Missing Values

Since all of the columns above are categorical columns, we will be filling them with the string 'unknown' so that we can maintain as much data as possible.

```
In [35]: wells_df.fillna('unknown', inplace=True)
```

```
In [36]: wells_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 57570 entries, 69572 to 26348
Data columns (total 40 columns):
#   Column                Non-Null Count  Dtype
---  -
0   status_group          57570 non-null  int64
1   amount_tsh            57570 non-null  float64
2   date_recorded         57570 non-null  object
3   funder                57570 non-null  object
4   gps_height            57570 non-null  int64
5   installer             57570 non-null  object
6   longitude             57570 non-null  float64
7   latitude              57570 non-null  float64
8   wpt_name              57570 non-null  object
9   num_private           57570 non-null  int64
10  basin                 57570 non-null  object
11  subvillage            57570 non-null  object
12  region                57570 non-null  object
13  region_code           57570 non-null  object
14  district_code         57570 non-null  object
15  lga                   57570 non-null  object
16  ward                  57570 non-null  object
17  population            57570 non-null  int64
18  public_meeting        57570 non-null  object
19  recorded_by           57570 non-null  object
20  scheme_management     57570 non-null  object
21  scheme_name           57570 non-null  object
22  permit               57570 non-null  object
23  construction_year     57570 non-null  int64
24  extraction_type       57570 non-null  object
25  extraction_type_group  57570 non-null  object
26  extraction_type_class  57570 non-null  object
27  management            57570 non-null  object
28  management_group      57570 non-null  object
29  payment               57570 non-null  object
30  payment_type          57570 non-null  object
31  water_quality         57570 non-null  object
32  quality_group         57570 non-null  object
33  quantity              57570 non-null  object
34  quantity_group        57570 non-null  object
35  source                57570 non-null  object
36  source_type           57570 non-null  object
37  source_class          57570 non-null  object
38  waterpoint_type       57570 non-null  object
39  waterpoint_type_group 57570 non-null  object
dtypes: float64(3), int64(5), object(32)
memory usage: 18.0+ MB
```

## Dropping Columns

```
In [37]: are_identical = wells_df['quantity'].equals(wells_df['quantity_group'])
print('Are they identical?', are_identical)
```

Are they identical? True

We are going to drop the following columns:

- 'recorded\_by'
  - All records were recorded by the same person. No unique data is provided.
- 'quantity\_group'
  - This column is identical to 'quantity'
- 'date\_recorded'
  - The date that the information was recorded does not hold valuable information for determining the status of the well.
- 'num\_private'
  - There is no description for this column inside of the data table's documentation.
  - <https://www.drivendata.org/competitions/7/pump-it-up-data-mining-the-water-table/page/25/>

```
In [38]: # Drop columns
wells_df.drop(['recorded_by', 'quantity_group', 'date_recorded', 'num_private'], axis=1, inplace=True)
```

## Baseline Feature Importance

### Baseline Feature Importance: Decision Tree

```
In [39]: baseline_dt_df = wells_df.copy()
```

```

In [40]: # Encode categorical variables using LabelEncoder
le = LabelEncoder()
for col in baseline_dt_df.select_dtypes(include='object').columns:
    baseline_dt_df[col] = le.fit_transform(baseline_dt_df[col])

In [41]: X_train, X_test, y_train, y_test = split_data(baseline_dt_df, 'status_group')

In [42]: # Train a simple decision tree classifier
dt = DecisionTreeClassifier(random_state=42)
dt.fit(X_train, y_train)

Out[42]: DecisionTreeClassifier(random_state=42)

In [43]: # Get feature importances
importances = dt.feature_importances_
feature_names = X_train.columns
feature_importance_df = pd.DataFrame({'Feature': feature_names, 'Importance': importances})

In [44]: # Sort by importance
feature_importance_df = feature_importance_df.sort_values(by='Importance', ascending=False)
feature_importance_df

```

Out[44]:

	Feature	Importance
29	quantity	0.143414
4	longitude	0.109715
5	latitude	0.095917
34	waterpoint_type_group	0.074964
6	wpt_name	0.072433
8	subvillage	0.063031
2	gps_height	0.047830
19	construction_year	0.038912
13	ward	0.037973
14	population	0.035885
1	funder	0.030737
3	installer	0.024856
0	amount_tsh	0.024518
12	lga	0.023856
17	scheme_name	0.022924
9	region	0.019552
22	extraction_type_class	0.014009
23	management	0.010207
20	extraction_type	0.009441
16	scheme_management	0.008787
18	permit	0.008369
25	payment	0.008288
11	district_code	0.008110
30	source	0.007853
31	source_type	0.007844
33	waterpoint_type	0.007246
7	basin	0.006857
21	extraction_type_group	0.006657
28	quality_group	0.005540
10	region_code	0.005316
26	payment_type	0.004768
27	water_quality	0.004146
24	management_group	0.004025
15	public_meeting	0.003425
32	source_class	0.002597

## Baseline Feature Importance: RFE

```
In [45]: # Initialize RFE
rfe = RFE(estimator=dt, n_features_to_select=10)
```

```
In [46]: # Fit RFE
rfe.fit(X_train, y_train)
```

```
Out[46]: RFE(estimator=DecisionTreeClassifier(random_state=42), n_features_to_select=10)
```

```
In [47]: # Get the selected features
selected_features = X_train.columns[rfe.support_]
print("Selected Features:")
print(selected_features)
```

Selected Features:

```
Index(['funder', 'gps_height', 'longitude', 'latitude', 'wpt_name',
      'subvillage', 'ward', 'construction_year', 'quantity',
      'waterpoint_type'],
      dtype='object')
```

```
In [48]: # Get feature ranking
ranking = rfe.ranking_
feature_ranking = pd.DataFrame({'Feature': X_train.columns, 'Ranking': ranking})

# Sort by ranking
feature_ranking = feature_ranking.sort_values(by='Ranking')
print("Features ranked by importance:")
print(feature_ranking)
```

Features ranked by importance:

	Feature	Ranking
1	funder	1
2	gps_height	1
4	longitude	1
5	latitude	1
6	wpt_name	1
29	quantity	1
8	subvillage	1
19	construction_year	1
13	ward	1
33	waterpoint_type	1
14	population	2
12	lga	3
0	amount_tsh	4
3	installer	5
22	extraction_type_class	6
9	region	7
17	scheme_name	8
23	management	9
30	source	10
25	payment	11
20	extraction_type	12
16	scheme_management	13
18	permit	14
28	quality_group	15
11	district_code	16
7	basin	17
21	extraction_type_group	18
31	source_type	19
10	region_code	20
26	payment_type	21
15	public_meeting	22
27	water_quality	23
24	management_group	24
34	waterpoint_type_group	25
32	source_class	26

## Baseline Feature Importance: Analysis

```
In [49]: # Merge feature importance and ranking DataFrames on feature names
baseline_feature_analysis_df = pd.merge(
    feature_importance_df, feature_ranking, on='Feature', suffixes=('_importance', '_ranking'))
```

Now that we have combined the feature selections together we can filter the dataframe so we can see the features that are less likely to contribute to our models performance.

We will want to look at any feature that matches both of the criteria listed below and consider dropping those features.

For RFE (Recursive Feature Elimination) features with higher ranks are considered less important by RFE so we will look at features with a rank below 20.

For the decision tree feature importance, we will look at features that have a gini impurity below .01.

```
In [50]: # Filter the DataFrame
features_to_drop = baseline_feature_analysis_df[
    (baseline_feature_analysis_df['Importance'] < 0.01) &
    (baseline_feature_analysis_df['Ranking'] >= 20)
][['Feature']]
```

```
In [51]: features_to_drop
```

```
Out[51]: 29      region_code
30      payment_type
31      water_quality
32      management_group
33      public_meeting
34      source_class
Name: Feature, dtype: object
```

```
In [52]: columns_to_drop = {}

for col in features_to_drop:
    columns_to_drop[col] = {
        'dtype' : wells_df[col].dtype,
        'unique_values' : wells_df[col].nunique()
    }

columns_to_drop
```

```
Out[52]: {'region_code': {'dtype': dtype('O'), 'unique_values': 27},
          'payment_type': {'dtype': dtype('O'), 'unique_values': 6},
          'water_quality': {'dtype': dtype('O'), 'unique_values': 8},
          'management_group': {'dtype': dtype('O'), 'unique_values': 4},
          'public_meeting': {'dtype': dtype('O'), 'unique_values': 3},
          'source_class': {'dtype': dtype('O'), 'unique_values': 3}}
```

### 'source\_class' and 'source\_type'

```
In [53]: wells_df['source_class'].value_counts()
```

```
Out[53]: groundwater    44201
surface                13103
unknown                 266
Name: source_class, dtype: int64
```

```
In [54]: wells_df['source_type'].value_counts()
```

```
Out[54]: spring                17006
shallow well                 15498
borehole                     11697
river/lake                    10236
rainwater harvesting          2218
dam                           649
unknown                       266
Name: source_type, dtype: int64
```

Even though this a low scoring column, the water source of the well is most likely a good indicator of the status of the well itself. We will want to retain as much information as we can that can lead us to determining the functional status of a well.

### 'payment\_type'

```
In [55]: wells_df['payment_type'].value_counts()
```

```
Out[55]: never pay            24380
per bucket                    8953
unknown                       8556
monthly                       8216
on failure                     3841
annually                       3624
Name: payment_type, dtype: int64
```

```
In [56]: wells_df['payment'].value_counts()
```

```
Out[56]: never pay                24380
        pay per bucket          8953
        unknown                 8556
        pay monthly             8216
        pay when scheme fails   3841
        pay annually            3624
        Name: payment, dtype: int64
```

This column is essentially the same as 'payment' as such we can drop this column.

## 'public\_meeting'

```
In [57]: wells_df['public_meeting'].value_counts()
```

```
Out[57]: True          49722
        False         4872
        unknown       2976
        Name: public_meeting, dtype: int64
```

Given its low scores and that there is no further information given in the documentation aside from True/False, we will drop this column.

## 'management\_group'

```
In [58]: wells_df['management_group'].value_counts()
```

```
Out[58]: user-group    50755
        commercial    3634
        parastatal     1691
        unknown        1490
        Name: management_group, dtype: int64
```

- How the waterpoint is managed?

This could indicate how different management practices that affect the functionality of the well. As such, we will want to keep this column.

## 'water\_quality'

```
In [59]: wells_df['water_quality'].value_counts()
```

```
Out[59]: soft          49413
        salty          4772
        unknown        1661
        milky           803
        coloured        479
        salty abandoned  228
        fluoride         199
        fluoride abandoned 15
        Name: water_quality, dtype: int64
```

```
In [60]: wells_df['quality_group'].value_counts()
```

```
Out[60]: good          49413
        salty          5000
        unknown        1661
        milky           803
        colored         479
        fluoride        214
        Name: quality_group, dtype: int64
```

While there is less crossover inside of water\_quality and quality\_group we will want to keep both.

The water quality of the well is most likely one of the best indicators of the status of the well.

We will want to retain as much information as we can that can lead us to determining the functional status of a well.

## Dropping columns

```
In [61]: wells_df = wells_df.drop(['payment_type', 'public_meeting'], axis=1)
```

## Columns with too many unique values

```
In [62]: wells_df.nunique()
```

```
Out[62]: status_group      2
amount_tsh      88
funder      1855
gps_height      2428
installer      1835
longitude      57497
latitude      57498
wpt_name      36708
basin      9
subvillage      18567
region      21
region_code      27
district_code      20
lga      124
ward      2033
population      1047
scheme_management      11
scheme_name      2537
permit      3
construction_year      55
extraction_type      18
extraction_type_group      13
extraction_type_class      7
management      11
management_group      4
payment      6
water_quality      8
quality_group      6
quantity      5
source      9
source_type      7
source_class      3
waterpoint_type      7
waterpoint_type_group      6
dtype: int64
```

Given the high amount of unique values that exist inside of various columns, we will bucket the values that fall under 1% of the columns total values.

We will only be bucketing values for columns that have over 30 unique values so that we maintain all of the unique values in columns that have smaller amount of unique values.

```
In [63]: too_many_unique_columns = []

for col in wells_df.columns:
    if wells_df[col].dtype == 'object' and wells_df[col].nunique() > 30:
        too_many_unique_columns.append(col)

too_many_unique_columns
```

```
Out[63]: ['funder', 'installer', 'wpt_name', 'subvillage', 'lga', 'ward', 'scheme_name']
```

```
In [64]: for col in too_many_unique_columns:
    if wells_df[col].dtype == 'object':
        freq = wells_df[col].value_counts(normalize=True)
        mask = wells_df[col].isin(freq[freq >= 0.01].index)
        wells_df.loc[~mask, col] = 'other'
```

```
In [65]: wells_df.nunique()
```

```
Out[65]: status_group      2
amount_tsh      88
funder          19
gps_height     2428
installer       15
longitude      57497
latitude       57498
wpt_name        4
basin           9
subvillage      1
region         21
region_code     27
district_code   20
lga            35
ward           1
population     1047
scheme_management 11
scheme_name      3
permit          3
construction_year 55
extraction_type  18
extraction_type_group 13
extraction_type_class 7
management     11
management_group 4
payment         6
water_quality   8
quality_group    6
quantity         5
source          9
source_type      7
source_class     3
waterpoint_type  7
waterpoint_type_group 6
dtype: int64
```

## Column Correlation

### Numeric Columns


```
In [66]: correlation_matrix = wells_df.corr()
```

```
In [67]: # Set up the matplotlib figure
plt.figure(figsize=(15, 10))

# Draw the heatmap with the correlation matrix
sns.heatmap(correlation_matrix, annot=True, fmt='.2f', cmap='coolwarm', linewidths=0.5)

# Add title and labels
plt.title('Correlation Matrix of Features')
plt.show()
```





While some columns border on being too correlated with each other ( $\sim 0.80$ ), none of columns exceed this threshold. If a column had multiple high correlations (either positive or negative) with other columns we would drop that column, but this is not the case.

## Scoring Metric

Inside of all the models listed below, we are using recall as our scoring metric.

In the target variable, we started with three classes:

- Functional
- Non-Functional
- Functional needs repair

We combined the non-functional and the functional needs repair classes due to the heavy class imbalance inside of the functional needs repair class. Since functional needs repair and non-functional wells could both be classified into a needs attention class we combined them. Our updated classes are:

- Does not need attention (for functional wells)
- Needs attention (for non-functional and functionally needs repair wells)

We chose to use recall as our primary scoring metric because it measures the proportion of wells that need attention that are correctly predicted as being in need of attention.

## Sampling Methods

```
In [68]: sampling_methods = ['oversample', 'undersample', 'smote']
```

Inside of each model we will be performing these three sampling methods to determine which method performs the best.

## Base Logistic Regression Models

```
In [69]: base_log_model_df = wells_df.copy()
```

```
In [70]: X_train, X_test, y_train, y_test = split_data(base_log_model_df, 'status_group')
```

```
In [71]: base_logistic_model = LogisticRegression(random_state=42, max_iter=2000)
```

```
In [72]: base_log_cv = CustomCrossValidator(base_logistic_model, X_train, y_train)
```

```
In [73]: base_log_cv.run_sampling_methods(sampling_methods, features=None, features_from='base', folds=5)
```

Running cross-validation with sampling method: oversample

Running cross-validation with sampling method: undersample

Running cross-validation with sampling method: smote

Log: {'LogisticRegression', 'base', 'oversample': 0.7043794944655426, ('LogisticRegression', 'base', 'undersample'): 0.7032904517477562, ('LogisticRegression', 'base', 'smote'): 0.7011659423295937}

## Base Decision Tree Models

```
In [74]: base_dt_model_df = wells_df.copy()
```

```
In [75]: X_train, X_test, y_train, y_test = split_data(base_dt_model_df, 'status_group')
```

```
In [76]: base_dt_model = DecisionTreeClassifier(random_state=42, max_depth=5, min_samples_split=20, min_samples_leaf=10)
```

```
In [77]: base_dt_cv = CustomCrossValidator(base_dt_model, X_train, y_train)
```

```
In [78]: base_dt_cv.run_sampling_methods(sampling_methods, features=None, features_from='base', folds=5)
```

Running cross-validation with sampling method: oversample

Running cross-validation with sampling method: undersample

Running cross-validation with sampling method: smote

Log: {'LogisticRegression', 'base', 'oversample': 0.7043794944655426, ('LogisticRegression', 'base', 'undersample'): 0.7032904517477562, ('LogisticRegression', 'base', 'smote'): 0.7011659423295937, ('DecisionTreeClassifier', 'base', 'oversample'): 0.5150329468026201, ('DecisionTreeClassifier', 'base', 'undersample'): 0.5310489564026828, ('DecisionTreeClassifier', 'base', 'smote'): 0.5393838624465799}

## L1 Penalty Logistic Regression: Feature Selection

```
In [79]: l1_penalty_log_regression_df = wells_df.copy()
```

### Feature Selection

```
In [80]: # Split the data
X_train, X_test, y_train, y_test = split_data(l1_penalty_log_regression_df, target='status_group')

# Scale numeric features
X_train_numeric_scaled, X_test_numeric_scaled = scale_numeric_features(X_train, X_test)

# Encode categorical features
X_train_encoded, X_test_encoded = encode_categorical_features(X_train, X_test)

# Combine numeric and categorical features
X_train_preprocessed = combine_features(X_train_numeric_scaled, X_train_encoded)
X_test_preprocessed = combine_features(X_test_numeric_scaled, X_test_encoded)
```

```
In [81]: # Instantiate the model
l1_penalty_log_regression_model = LogisticRegression(random_state=42, penalty='l1', solver='liblinear')
```

```
In [82]: # Fit the model
l1_penalty_log_regression_model.fit(X_train_preprocessed, y_train)
```

```
Out[82]: LogisticRegression(penalty='l1', random_state=42, solver='liblinear')
```

```
In [83]: # Extract feature importance (non-zero coefficients)
importance = l1_penalty_log_regression_model.coef_[0]

l1_penalty_selected_features = X_train_preprocessed.columns[importance != 0]
```

```
In [84]: print(f"Starting Features: {X_train_preprocessed.shape[1]}")
print()
print(f"Selected Features: {len(l1_penalty_selected_features)}")
```

Starting Features: 284

Selected Features: 224

The L1 penalty on the Logistic Regression set the coefficient of 60 features to 0 indicating that these features are not considered important for predicting the target variable.

## L1 Penalty Features: Logistic Regression

```
In [85]: l1_features_l2_logistic_model_df = wells_df.copy()

In [86]: # Split the data
X_train, X_test, y_train, y_test = split_data(l1_features_l2_logistic_model_df, target='status_group')

In [87]: # Instantiate the model
l1_features_l2_logistic_model = LogisticRegression(random_state=42, max_iter=2000)

In [88]: l1_features_l2_logistic_cv = CustomCrossValidator(
    model=l1_features_l2_logistic_model,
    X=X_train,
    y=y_train,
)

In [89]: l1_features_l2_logistic_cv.run_sampling_methods(
    sampling_methods,
    features=l1_penalty_selected_features,
    features_from='L1_Penalty'
)
```

Running cross-validation with sampling method: oversample

Running cross-validation with sampling method: undersample

Running cross-validation with sampling method: smote

Log: {'LogisticRegression', 'base', 'oversample': 0.7043794944655426, ('LogisticRegression', 'base', 'undersample'): 0.7032904517477562, ('LogisticRegression', 'base', 'smote'): 0.7011659423295937, ('DecisionTreeClassifier', 'base', 'oversample'): 0.5150329468026201, ('DecisionTreeClassifier', 'base', 'undersample'): 0.5310489564026828, ('DecisionTreeClassifier', 'base', 'smote'): 0.5393838624465799, ('LogisticRegression', 'L1\_Penalty', 'oversample'): 0.7048153281712818, ('LogisticRegression', 'L1\_Penalty', 'undersample'): 0.7038352475891534, ('LogisticRegression', 'L1\_Penalty', 'smote'): 0.7014929177579201}

## L1 Penalty Features: Decision Tree

```
In [90]: l1_features_dt_model_df = wells_df.copy()

X_train, X_test, y_train, y_test = split_data(l1_features_dt_model_df, 'status_group')

In [91]: l1_features_dt_model = DecisionTreeClassifier(random_state=42, max_depth=5, min_samples_split=20, min_samples_leaf=10)

In [92]: l1_features_dt_cv = CustomCrossValidator(
    model=l1_features_dt_model,
    X=X_train,
    y=y_train,
)

In [93]: l1_features_dt_cv.run_sampling_methods(
    sampling_methods,
    features=l1_penalty_selected_features,
    folds=5,
    features_from='L1_Penalty'
)
```

Running cross-validation with sampling method: oversample

Running cross-validation with sampling method: undersample

Running cross-validation with sampling method: smote

Log: {'LogisticRegression', 'base', 'oversample': 0.7043794944655426, ('LogisticRegression', 'base', 'undersample'): 0.7032904517477562, ('LogisticRegression', 'base', 'smote'): 0.7011659423295937, ('DecisionTreeClassifier', 'base', 'oversample'): 0.5150329468026201, ('DecisionTreeClassifier', 'base', 'undersample'): 0.5310489564026828, ('DecisionTreeClassifier', 'base', 'smote'): 0.5393838624465799, ('LogisticRegression', 'L1\_Penalty', 'oversample'): 0.7048153281712818, ('LogisticRegression', 'L1\_Penalty', 'undersample'): 0.7038352475891534, ('LogisticRegression', 'L1\_Penalty', 'smote'): 0.7014929177579201, ('DecisionTreeClassifier', 'L1\_Penalty', 'oversample'): 0.5151418792644937, ('DecisionTreeClassifier', 'L1\_Penalty', 'undersample'): 0.5311034374705116, ('DecisionTreeClassifier', 'L1\_Penalty', 'smote'): 0.539274915147814}

# VIF: Feature Selection

```
In [94]: vif_df = wells_df.copy()
```

## Feature Selection

```
In [95]: # Split the data
X_train, X_test, y_train, y_test = split_data(vif_df, target='status_group')

# Scale numeric features
X_train_numeric_scaled, X_test_numeric_scaled = scale_numeric_features(X_train, X_test)

# Encode categorical features
X_train_encoded, X_test_encoded = encode_categorical_features(X_train, X_test)

# Combine numeric and categorical features
X_train_preprocessed = combine_features(X_train_numeric_scaled, X_train_encoded)
X_test_preprocessed = combine_features(X_test_numeric_scaled, X_test_encoded)
```

```
In [96]: # Remove features with variance below 1%
selector = VarianceThreshold(threshold=0.01)
X_train_reduced_df = selector.fit_transform(X_train_preprocessed)

# Get the selected feature names
selected_features = X_train_preprocessed.columns[selector.get_support()]

# Create a new DataFrame with selected features
X_train_reduced_df = pd.DataFrame(X_train_reduced_df, columns=selected_features)
```

```
In [97]: X_train_reduced_df.shape
```

```
Out[97]: (40299, 234)
```

```
In [98]: # Function to calculate VIF
def calculate_vif(df):
    vif_data = pd.DataFrame()
    vif_data['feature'] = df.columns
    vif_data['VIF'] = [variance_inflation_factor(df.values, i) for i in range(df.shape[1])]
    return vif_data

# Apply VIF calculation on reduced DataFrame
vif_data_reduced = calculate_vif(X_train_reduced_df)

vif_data_reduced
```

```
/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages/statsmodels/stats/outliers_influence.py:193: RuntimeWarning: divide by zero encountered in double_scalars
  vif = 1. / (1. - r_squared_i)
```

```
Out[98]:
```

	feature	VIF
0	amount_tsh	1.326252
1	gps_height	11.058491
2	longitude	90.910726
3	latitude	94.372529
4	population	1.224808
...	...	...
229	waterpoint_type_unknown	inf
230	waterpoint_type_group_communal standpipe	inf
231	waterpoint_type_group_hand pump	inf
232	waterpoint_type_group_improved spring	inf
233	waterpoint_type_group_unknown	inf

234 rows × 2 columns

```
In [99]: vif_under_10 = vif_data_reduced[vif_data_reduced['VIF'] < 10]
vif_under_10
```

Out[99]:

	feature	VIF
0	amount_tsh	1.326252
4	population	1.224808
7	funder_dhv	3.541435
8	funder_district council	3.381726
11	funder_hesawa	6.251303
12	funder_kkkt	6.220374
13	funder_norad	3.205758
15	funder_private individual	3.597767
16	funder_rwssp	4.516637
17	funder_tasaf	2.994374
18	funder_tcrs	4.749076
19	funder_unicef	3.548683
21	funder_water	2.764913
22	funder_world bank	4.081632
23	funder_world vision	5.146454
99	district_code_33	4.070209
102	district_code_53	1.898749
106	lga_arusha rural	9.185187
107	lga_bagamoyo	8.076241
109	lga_iringa rural	4.935398
110	lga_kahama	4.505269
111	lga_karagwe	4.632777
115	lga_kilombero	7.291015
116	lga_kilosa	7.363698
117	lga_kwimba	4.234895
118	lga_kyela	6.953799
119	lga_lushoto	4.525396
120	lga_makete	4.769707
121	lga_maswa	3.617241
122	lga_mbarali	5.927934
123	lga_mbinga	6.360571
124	lga_mbozi	9.059500
125	lga_meru	9.479916
126	lga_moshi rural	7.421927
127	lga_mpanda	5.476681
128	lga_mvomero	5.587916
129	lga_namtumbo	6.492646
133	lga_rombo	7.111824
135	lga_same	6.307670
136	lga_serengeti	5.374103
137	lga_singida rural	9.735145
138	lga_songea rural	6.203831
139	lga_ulanga	5.873941

```
In [100]: vif_selected_features = list(vif_under_10['feature'])
          len(vif_selected_features)
```

Out[100]: 43

## VIF Features: Logistic Regression

```
In [101.. vif_logistic_model_df = wells_df.copy()

X_train, X_test, y_train, y_test = split_data(vif_logistic_model_df, 'status_group')
```

```
In [102.. vif_logistic_model = LogisticRegression(random_state=42, max_iter=2000)
```

```
In [103.. vif_logistic_cv = CustomCrossValidator(
    model=vif_logistic_model,
    X=X_train,
    y=y_train,
)
```

```
In [104.. vif_logistic_cv.run_sampling_methods(
    sampling_methods,
    features=vif_selected_features,
    folds=5,
    features_from='VIF'
)
```

Running cross-validation with sampling method: oversample

Running cross-validation with sampling method: undersample

Running cross-validation with sampling method: smote

```
Log: {'LogisticRegression', 'base', 'oversample'): 0.7043794944655426, ('LogisticRegression', 'base', 'undersam
ple'): 0.7032904517477562, ('LogisticRegression', 'base', 'smote'): 0.7011659423295937, ('DecisionTreeClassifier
', 'base', 'oversample'): 0.5150329468026201, ('DecisionTreeClassifier', 'base', 'undersample'): 0.5310489564026
828, ('DecisionTreeClassifier', 'base', 'smote'): 0.5393838624465799, ('LogisticRegression', 'L1_Penalty', 'over
sample'): 0.7048153281712818, ('LogisticRegression', 'L1_Penalty', 'undersample'): 0.7038352475891534, ('Logisti
cRegression', 'L1_Penalty', 'smote'): 0.7014929177579201, ('DecisionTreeClassifier', 'L1_Penalty', 'oversample')
: 0.5151418792644937, ('DecisionTreeClassifier', 'L1_Penalty', 'undersample'): 0.5311034374705116, ('DecisionTre
eClassifier', 'L1_Penalty', 'smote'): 0.539274915147814, ('LogisticRegression', 'VIF', 'oversample'): 0.68063043
73500361, ('LogisticRegression', 'VIF', 'undersample'): 0.6832991788076954, ('LogisticRegression', 'VIF', 'smote
'): 0.6713702433665739}
```

## VIF Features: Decision Tree

```
In [105.. vif_dt_model_df = wells_df.copy()

X_train, X_test, y_train, y_test = split_data(vif_dt_model_df, 'status_group')
```

```
In [106.. vif_dt_model = DecisionTreeClassifier(random_state=42, max_depth=5, min_samples_split=20, min_samples_leaf=10)
```

```
In [107.. vif_dt_cv = CustomCrossValidator(
    model=vif_dt_model,
    X=X_train,
    y=y_train,
)
```

```
In [108.. vif_dt_cv.run_sampling_methods(
    sampling_methods,
    features=vif_selected_features,
    folds=5,
    features_from='VIF'
)
```

Running cross-validation with sampling method: oversample

Running cross-validation with sampling method: undersample

Running cross-validation with sampling method: smote

```
Log: {'LogisticRegression', 'base', 'oversample'): 0.7043794944655426, ('LogisticRegression', 'base', 'undersam
ple'): 0.7032904517477562, ('LogisticRegression', 'base', 'smote'): 0.7011659423295937, ('DecisionTreeClassifier
', 'base', 'oversample'): 0.5150329468026201, ('DecisionTreeClassifier', 'base', 'undersample'): 0.5310489564026
828, ('DecisionTreeClassifier', 'base', 'smote'): 0.5393838624465799, ('LogisticRegression', 'L1_Penalty', 'over
sample'): 0.7048153281712818, ('LogisticRegression', 'L1_Penalty', 'undersample'): 0.7038352475891534, ('Logisti
cRegression', 'L1_Penalty', 'smote'): 0.7014929177579201, ('DecisionTreeClassifier', 'L1_Penalty', 'oversample')
: 0.5151418792644937, ('DecisionTreeClassifier', 'L1_Penalty', 'undersample'): 0.5311034374705116, ('DecisionTre
eClassifier', 'L1_Penalty', 'smote'): 0.539274915147814, ('LogisticRegression', 'VIF', 'oversample'): 0.68063043
73500361, ('LogisticRegression', 'VIF', 'undersample'): 0.6832991788076954, ('LogisticRegression', 'VIF', 'smote
'): 0.6713702433665739, ('DecisionTreeClassifier', 'VIF', 'oversample'): 0.7701797459805376, ('DecisionTreeClass
ifier', 'VIF', 'undersample'): 0.767510559416115, ('DecisionTreeClassifier', 'VIF', 'smote'): 0.768382701608141}
```

## Current Model Results

```
In [109.. def get_log_results_df():
    # Extracting the log dictionary
```

```

log_dict = vif_dt_cv.log

# Initialize an empty list to store the rows
log_list = []

# Iterate over the log dictionary and append each entry as a tuple to the list
for key, value in log_dict.items():
    model, feature_selection, sampling = key
    score = value
    log_list.append((model, feature_selection, sampling, score))

# Convert the list to a DataFrame
log_df = pd.DataFrame(log_list, columns=['Model', 'Feature Selection', 'Sampling Technique', 'Recall-Score'])

return log_df

```

```

In [110]: log_df = get_log_results_df()
log_df.sort_values(by='Recall-Score', ascending=False)

```

```

Out[110]:

```

	Model	Feature Selection	Sampling Technique	Recall-Score
15	DecisionTreeClassifier	VIF	oversample	0.770180
17	DecisionTreeClassifier	VIF	smote	0.768383
16	DecisionTreeClassifier	VIF	undersample	0.767511
6	LogisticRegression	L1_Penalty	oversample	0.704815
0	LogisticRegression	base	oversample	0.704379
7	LogisticRegression	L1_Penalty	undersample	0.703835
1	LogisticRegression	base	undersample	0.703290
8	LogisticRegression	L1_Penalty	smote	0.701493
2	LogisticRegression	base	smote	0.701166
13	LogisticRegression	VIF	undersample	0.683299
12	LogisticRegression	VIF	oversample	0.680630
14	LogisticRegression	VIF	smote	0.671370
5	DecisionTreeClassifier	base	smote	0.539384
11	DecisionTreeClassifier	L1_Penalty	smote	0.539275
10	DecisionTreeClassifier	L1_Penalty	undersample	0.531103
4	DecisionTreeClassifier	base	undersample	0.531049
9	DecisionTreeClassifier	L1_Penalty	oversample	0.515142
3	DecisionTreeClassifier	base	oversample	0.515033

Given the success of the base logistic regression model, we want to perform some dimensionality reduction to see if we can get higher recall scores from our model.

We will perform dimensionality reduction on columns that share similar values. We kept these columns in previously because we believed that while they were similar, they did have differences and those differences could lead to a better model performance.

## Dimensionality Reduction

Various columns contain similar information. While these columns are not a equal in every regard, they are similar enough as to where our models may perform better if we removed these features. We had originally kept these features as we thought having the extra data points inside of these features would be helpful for classification.

If the columns contain over 80% of shared row values, we will drop a column.

We will be comparing the following features:

- 'scheme\_management' and 'management'
- 'region' and 'region\_code'
- 'extraction\_type' and 'extraction\_type\_group'
- 'water\_quality' and 'quality\_group'
- 'source' and 'source\_type'
- 'waterpoint\_type' and 'waterpoint\_type\_group'

```

In [111]: def get_shared_values_percentage(df, pairs):
"""
Calculates the percentage of shared values between pairs of columns in a DataFrame.

```

```

Return a dictionary with the shared values percentage.
---
Input:
df: Pandas DataFrame
pairs : A list of tuples
---
Output:
dict: A dictionary
"""
result = {}

for column1, column2 in pairs:
    # Create a new DataFrame with only the two specified columns
    df_subset = df[[column1, column2]].copy()

    # Create a new column that indicates if the values in the two columns are the same
    df_subset['is_same'] = df_subset[column1] == df_subset[column2]

    # Count the number of shared values (where the values are the same in both columns)
    shared_values_count = df_subset['is_same'].sum()

    # Calculate the percentage of shared values
    shared_values_percentage = (shared_values_count / len(df_subset)) * 100

    # Store the result in the dictionary
    result[(column1, column2)] = shared_values_percentage

return result

```

```

In [112]: # Define the list of column pairs to compare
pairs = [
    ('scheme_management', 'management'),
    ('extraction_type', 'extraction_type_group'),
    ('water_quality', 'quality_group'),
    ('source', 'source_type'),
    ('waterpoint_type', 'waterpoint_type_group')
]

```

```

In [113]: # Calculate shared values percentages for each pair
results = get_shared_values_percentage(wells_df, pairs)
results

```

```

Out[113]: {('scheme_management', 'management'): 83.8909154073302,
('extraction_type', 'extraction_type_group'): 96.64929650859823,
('water_quality', 'quality_group'): 12.914712523883967,
('source', 'source_type'): 61.902032308494,
('waterpoint_type', 'waterpoint_type_group'): 89.65085982282439}

```

Given the results above, we will take a closer look at the following pairs:

- 'scheme\_management' and 'management'
- 'extraction\_type' and 'extraction\_type\_group'
- 'waterpoint\_type' and 'waterpoint\_type\_group'

We will also take a look at 'region' and 'region\_code' separately as while the values are not the same, they contain similar information.

## 'scheme\_management' and 'management'

```

In [114]: wells_df['scheme_management'].value_counts()

```

```

Out[114]: vwc          36140
unknown      4516
wug          4248
water authority  3150
wua          2872
water board  2747
parastatal   1605
private operator  1063
company      1060
swc           97
trust         72
Name: scheme_management, dtype: int64

```

```

In [115]: wells_df['management'].value_counts()

```



```
Out[115...] vwc                39743
            wug                5554
            water board       2932
            wua                2526
            private operator   1969
            parastatal         1691
            unknown           1391
            water authority     902
            company            685
            school             99
            trust              78
            Name: management, dtype: int64
```

Since management has fewer unknown values, we will drop the 'scheme\_management' column from the DataFrame.

## 'extraction\_type' and 'extraction\_type\_group'

```
In [116...] wells_df['extraction_type'].value_counts()
```

```
Out[116...] gravity                26687
            nira/tanira             7361
            unknown                 6160
            submersible             4687
            swm 80                  3447
            mono                   2815
            india mark ii           2283
            afridev                 1659
            ksb                    1354
            rope pump               451
            swm 81                  229
            windmill               117
            india mark iii           91
            cemo                    90
            play pump               85
            climax                  32
            walimi                  20
            mkulima/shinyanga        2
            Name: extraction_type, dtype: int64
```

```
In [117...] wells_df['extraction_type_group'].value_counts()
```

```
Out[117...] gravity                26687
            nira/tanira             7361
            unknown                 6160
            submersible             6041
            swm 80                  3447
            mono                   2815
            india mark ii           2283
            afridev                 1659
            rope pump               451
            other handpump           336
            other motorpump          122
            wind-powered            117
            india mark iii           91
            Name: extraction_type_group, dtype: int64
```

Since the 'extraction\_type' column is more specific than 'extraction\_type\_group', we will be dropping 'extraction\_type\_group' from the DataFrame.

## 'waterpoint\_type' and 'waterpoint\_type\_group'

```
In [118...] wells_df['waterpoint_type'].value_counts()
```

```
Out[118...] communal standpipe      28360
            hand pump               16179
            unknown                 6167
            communal standpipe multiple 5958
            improved spring           783
            cattle trough            116
            dam                       7
            Name: waterpoint_type, dtype: int64
```

```
In [119...] wells_df['waterpoint_type_group'].value_counts()
```

```
Out[119.. communal standpipe      34318
           hand pump              16179
           unknown                6167
           improved spring         783
           cattle trough           116
           dam                     7
           Name: waterpoint_type_group, dtype: int64
```

Since the 'waterpoint\_type' column is more specific than 'waterpoint\_type\_group', we will be dropping 'waterpoint\_type\_group' from the DataFrame.

## 'region' and 'region\_code'

```
In [120.. len(wells_df['region'].value_counts())
```

```
Out[120.. 21
```

```
In [121.. len(wells_df['region_code'].value_counts())
```

```
Out[121.. 27
```

Since the 'region\_code' contains more values than the region, we will be dropping 'region' from the DataFrame.

## Dropping columns

```
In [122.. dimensionality_reduction_df = wells_df.copy()
```

```
In [123.. dimensionality_reduction_df.drop(['scheme_management', 'extraction_type_group', 'waterpoint_type_group', 'region'
```

```
In [124.. dimensionality_reduction_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 57570 entries, 69572 to 26348
Data columns (total 30 columns):
#   Column                Non-Null Count  Dtype
---  -
0   status_group          57570 non-null  int64
1   amount_tsh            57570 non-null  float64
2   funder                57570 non-null  object
3   gps_height            57570 non-null  int64
4   installer             57570 non-null  object
5   longitude             57570 non-null  float64
6   latitude              57570 non-null  float64
7   wpt_name              57570 non-null  object
8   basin                57570 non-null  object
9   subvillage            57570 non-null  object
10  region_code           57570 non-null  object
11  district_code         57570 non-null  object
12  lga                   57570 non-null  object
13  ward                  57570 non-null  object
14  population            57570 non-null  int64
15  scheme_name           57570 non-null  object
16  permit                57570 non-null  object
17  construction_year     57570 non-null  int64
18  extraction_type        57570 non-null  object
19  extraction_type_class  57570 non-null  object
20  management             57570 non-null  object
21  management_group      57570 non-null  object
22  payment               57570 non-null  object
23  water_quality          57570 non-null  object
24  quality_group         57570 non-null  object
25  quantity              57570 non-null  object
26  source                57570 non-null  object
27  source_type           57570 non-null  object
28  source_class          57570 non-null  object
29  waterpoint_type       57570 non-null  object
dtypes: float64(3), int64(4), object(23)
memory usage: 13.6+ MB
```

## Dimensionality Reduction: Logistic Model

```
In [125.. dim_reduct_logistic_model_df = dimensionality_reduction_df.copy()
```

```
In [126.. X_train, X_test, y_train, y_test = split_data(dim_reduct_logistic_model_df, 'status_group')
```

```
In [127.. dim_reduct_logistic_model = LogisticRegression(random_state=42, max_iter=2000)
```

```
In [128.. dim_reduct_log_cv = CustomCrossValidator(dim_reduct_logistic_model, X_train, y_train)
```

```
In [129.. dim_reduct_log_cv.run_sampling_methods(sampling_methods, features=None, features_from='dimensionality_reduction')
```

Running cross-validation with sampling method: oversample

Running cross-validation with sampling method: undersample

Running cross-validation with sampling method: smote

Log: {'LogisticRegression', 'base', 'oversample'): 0.7043794944655426, ('LogisticRegression', 'base', 'undersample'): 0.7032904517477562, ('LogisticRegression', 'base', 'smote'): 0.7011659423295937, ('DecisionTreeClassifier', 'base', 'oversample'): 0.5150329468026201, ('DecisionTreeClassifier', 'base', 'undersample'): 0.5310489564026828, ('DecisionTreeClassifier', 'base', 'smote'): 0.5393838624465799, ('LogisticRegression', 'L1\_Penalty', 'oversample'): 0.7048153281712818, ('LogisticRegression', 'L1\_Penalty', 'undersample'): 0.7038352475891534, ('LogisticRegression', 'L1\_Penalty', 'smote'): 0.7014929177579201, ('DecisionTreeClassifier', 'L1\_Penalty', 'oversample'): 0.5151418792644937, ('DecisionTreeClassifier', 'L1\_Penalty', 'undersample'): 0.5311034374705116, ('DecisionTreeClassifier', 'L1\_Penalty', 'smote'): 0.539274915147814, ('LogisticRegression', 'VIF', 'oversample'): 0.6806304373500361, ('LogisticRegression', 'VIF', 'undersample'): 0.6832991788076954, ('LogisticRegression', 'VIF', 'smote'): 0.6713702433665739, ('DecisionTreeClassifier', 'VIF', 'oversample'): 0.7701797459805376, ('DecisionTreeClassifier', 'VIF', 'undersample'): 0.767510559416115, ('DecisionTreeClassifier', 'VIF', 'smote'): 0.768382701608141, ('LogisticRegression', 'dimensionality\_reduction', 'oversample'): 0.7030721120434614, ('LogisticRegression', 'dimensionality\_reduction', 'undersample'): 0.7016560790604568, ('LogisticRegression', 'dimensionality\_reduction', 'smote'): 0.6995859765256629}

## Dimensionality Reduction: Decision Tree

```
In [130.. dim_reduct_dt_model_df = dimensionality_reduction_df.copy()
```

```
In [131.. X_train, X_test, y_train, y_test = split_data(dim_reduct_dt_model_df, 'status_group')
```

```
In [132.. dim_reduct_dt_model = DecisionTreeClassifier(random_state=42, max_depth=5, min_samples_split=20, min_samples_leaf=10)
```

```
In [133.. dim_reduct_dt_cv = CustomCrossValidator(dim_reduct_dt_model, X_train, y_train)
```

```
In [134.. dim_reduct_dt_cv.run_sampling_methods(sampling_methods, features=None, features_from='dimensionality_reduction')
```

Running cross-validation with sampling method: oversample

Running cross-validation with sampling method: undersample

Running cross-validation with sampling method: smote

Log: {'LogisticRegression', 'base', 'oversample'): 0.7043794944655426, ('LogisticRegression', 'base', 'undersample'): 0.7032904517477562, ('LogisticRegression', 'base', 'smote'): 0.7011659423295937, ('DecisionTreeClassifier', 'base', 'oversample'): 0.5150329468026201, ('DecisionTreeClassifier', 'base', 'undersample'): 0.5310489564026828, ('DecisionTreeClassifier', 'base', 'smote'): 0.5393838624465799, ('LogisticRegression', 'L1\_Penalty', 'oversample'): 0.7048153281712818, ('LogisticRegression', 'L1\_Penalty', 'undersample'): 0.7038352475891534, ('LogisticRegression', 'L1\_Penalty', 'smote'): 0.7014929177579201, ('DecisionTreeClassifier', 'L1\_Penalty', 'oversample'): 0.5151418792644937, ('DecisionTreeClassifier', 'L1\_Penalty', 'undersample'): 0.5311034374705116, ('DecisionTreeClassifier', 'L1\_Penalty', 'smote'): 0.539274915147814, ('LogisticRegression', 'VIF', 'oversample'): 0.6806304373500361, ('LogisticRegression', 'VIF', 'undersample'): 0.6832991788076954, ('LogisticRegression', 'VIF', 'smote'): 0.6713702433665739, ('DecisionTreeClassifier', 'VIF', 'oversample'): 0.7701797459805376, ('DecisionTreeClassifier', 'VIF', 'undersample'): 0.767510559416115, ('DecisionTreeClassifier', 'VIF', 'smote'): 0.768382701608141, ('LogisticRegression', 'dimensionality\_reduction', 'oversample'): 0.7030721120434614, ('LogisticRegression', 'dimensionality\_reduction', 'undersample'): 0.7016560790604568, ('LogisticRegression', 'dimensionality\_reduction', 'smote'): 0.6995859765256629, ('DecisionTreeClassifier', 'dimensionality\_reduction', 'oversample'): 0.5150329468026201, ('DecisionTreeClassifier', 'dimensionality\_reduction', 'undersample'): 0.5311034374705116, ('DecisionTreeClassifier', 'dimensionality\_reduction', 'smote'): 0.5393293813787509}

## Review Model Results

```
In [135.. log_df = get_log_results_df()
log_df = log_df.sort_values(by='Recall-Score', ascending=False)
log_df
```

Out [135..

	Model	Feature Selection	Sampling Technique	Recall-Score
15	DecisionTreeClassifier	VIF	oversample	0.770180
17	DecisionTreeClassifier	VIF	smote	0.768383
16	DecisionTreeClassifier	VIF	undersample	0.767511
6	LogisticRegression	L1_Penalty	oversample	0.704815
0	LogisticRegression	base	oversample	0.704379
7	LogisticRegression	L1_Penalty	undersample	0.703835
1	LogisticRegression	base	undersample	0.703290
18	LogisticRegression	dimensionality_reduction	oversample	0.703072
19	LogisticRegression	dimensionality_reduction	undersample	0.701656
8	LogisticRegression	L1_Penalty	smote	0.701493
2	LogisticRegression	base	smote	0.701166
20	LogisticRegression	dimensionality_reduction	smote	0.699586
13	LogisticRegression	VIF	undersample	0.683299
12	LogisticRegression	VIF	oversample	0.680630
14	LogisticRegression	VIF	smote	0.671370
5	DecisionTreeClassifier	base	smote	0.539384
23	DecisionTreeClassifier	dimensionality_reduction	smote	0.539329
11	DecisionTreeClassifier	L1_Penalty	smote	0.539275
10	DecisionTreeClassifier	L1_Penalty	undersample	0.531103
22	DecisionTreeClassifier	dimensionality_reduction	undersample	0.531103
4	DecisionTreeClassifier	base	undersample	0.531049
9	DecisionTreeClassifier	L1_Penalty	oversample	0.515142
3	DecisionTreeClassifier	base	oversample	0.515033
21	DecisionTreeClassifier	dimensionality_reduction	oversample	0.515033

Our best models are the VIF Decision Tree models.

We will hyper-parameter tune all of these models to determine which model performs best of the validation fold.

Then we will use that model on the true hold out test.

## Hyperparameter Tuning

In [136..

```
decision_tree_param_grid = {
    'criterion': ['gini', 'entropy'],
    'max_depth': [5, 10, 15],
    'min_samples_split': [5, 10, 15],
    'min_samples_leaf': [4, 5, 6],
    'max_features': [None, 'sqrt', 'log2']
}
```

## VIF Oversampling Decision Tree Tuning

In [137..

```
tuned_vif_oversample_dt_model_df = wells_df.copy()

X_train, X_test, y_train, y_test = split_data(tuned_vif_oversample_dt_model_df, 'status_group')
```

In [138..

```
tuned_vif_oversample_dt_model = DecisionTreeClassifier(random_state=42)
```

In [139..

```
# Create an instance of CustomCrossValidator
tuned_vif_oversample_dt_cv = CustomCrossValidator(
    model=tuned_vif_oversample_dt_model,
    X=X_train,
    y=y_train,
    sampling_method='oversample'
)
```

In [140..

```
vif_oversample_dt_grid_search = tuned_vif_oversample_dt_cv.run_grid_search(
    param_grid=decision_tree_param_grid,
    features=vif_selected_features,
```

```
cv=StratifiedKFold(n_splits=5)
)
```

Fitting 5 folds for each of 162 candidates, totalling 810 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 34 tasks | elapsed: 5.7s
[Parallel(n_jobs=-1)]: Done 184 tasks | elapsed: 15.7s
[Parallel(n_jobs=-1)]: Done 434 tasks | elapsed: 30.2s
[Parallel(n_jobs=-1)]: Done 784 tasks | elapsed: 55.7s
[Parallel(n_jobs=-1)]: Done 810 out of 810 | elapsed: 56.7s finished
```

Fitting 5 folds for each of 162 candidates, totalling 810 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 34 tasks | elapsed: 3.0s
[Parallel(n_jobs=-1)]: Done 184 tasks | elapsed: 11.3s
[Parallel(n_jobs=-1)]: Done 434 tasks | elapsed: 28.0s
[Parallel(n_jobs=-1)]: Done 784 tasks | elapsed: 56.7s
[Parallel(n_jobs=-1)]: Done 810 out of 810 | elapsed: 58.0s finished
```

Fitting 5 folds for each of 162 candidates, totalling 810 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 34 tasks | elapsed: 3.5s
[Parallel(n_jobs=-1)]: Done 184 tasks | elapsed: 14.2s
[Parallel(n_jobs=-1)]: Done 434 tasks | elapsed: 30.7s
[Parallel(n_jobs=-1)]: Done 784 tasks | elapsed: 56.4s
[Parallel(n_jobs=-1)]: Done 810 out of 810 | elapsed: 57.5s finished
```

Fitting 5 folds for each of 162 candidates, totalling 810 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 34 tasks | elapsed: 3.3s
[Parallel(n_jobs=-1)]: Done 184 tasks | elapsed: 12.9s
[Parallel(n_jobs=-1)]: Done 434 tasks | elapsed: 30.9s
[Parallel(n_jobs=-1)]: Done 784 tasks | elapsed: 59.1s
[Parallel(n_jobs=-1)]: Done 810 out of 810 | elapsed: 1.0min finished
```

Fitting 5 folds for each of 162 candidates, totalling 810 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 34 tasks | elapsed: 3.3s
[Parallel(n_jobs=-1)]: Done 184 tasks | elapsed: 12.0s
[Parallel(n_jobs=-1)]: Done 434 tasks | elapsed: 25.9s
[Parallel(n_jobs=-1)]: Done 784 tasks | elapsed: 48.8s
[Parallel(n_jobs=-1)]: Done 810 out of 810 | elapsed: 49.9s finished
```

Mean Recall Score: 0.7284020548502097

Best parameters: {'criterion': 'gini', 'max\_depth': 5, 'max\_features': 'log2', 'min\_samples\_leaf': 4, 'min\_samples\_split': 5}

Best score: 0.7414135468421996

```
In [141]: vif_oversample_dt_grid_search_best_params = vif_oversample_dt_grid_search.best_params_
```

## VIF Undersampling Decision Tree Tuning

```
In [142]: tuned_vif_undersample_dt_model_df = wells_df.copy()
```

```
X_train, X_test, y_train, y_test = split_data(tuned_vif_undersample_dt_model_df, 'status_group')
```

```
In [143]: tuned_vif_undersample_dt_model = DecisionTreeClassifier(random_state=42)
```

```
In [144]: # Create an instance of CustomCrossValidator
tuned_vif_undersample_dt_cv = CustomCrossValidator(
    model=tuned_vif_undersample_dt_model,
    X=X_train,
    y=y_train,
    sampling_method='undersample'
)
```

```
In [145]: vif_undersample_dt_grid_search = tuned_vif_undersample_dt_cv.run_grid_search(
    param_grid=decision_tree_param_grid,
    features=vif_selected_features,
    cv=StratifiedKFold(n_splits=5)
)
```

Fitting 5 folds for each of 162 candidates, totalling 810 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 34 tasks | elapsed: 2.8s
[Parallel(n_jobs=-1)]: Done 184 tasks | elapsed: 10.0s
[Parallel(n_jobs=-1)]: Done 434 tasks | elapsed: 22.6s
[Parallel(n_jobs=-1)]: Done 784 tasks | elapsed: 40.3s
[Parallel(n_jobs=-1)]: Done 810 out of 810 | elapsed: 41.0s finished
```

Fitting 5 folds for each of 162 candidates, totalling 810 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed: 3.2s
[Parallel(n_jobs=-1)]: Done 184 tasks     | elapsed: 15.4s
[Parallel(n_jobs=-1)]: Done 434 tasks     | elapsed: 29.2s
```

```
[Parallel(n_jobs=-1)]: Done 784 tasks     | elapsed: 44.4s
[Parallel(n_jobs=-1)]: Done 810 out of 810 | elapsed: 44.9s finished
```

Fitting 5 folds for each of 162 candidates, totalling 810 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
```

```
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed: 2.3s
[Parallel(n_jobs=-1)]: Done 184 tasks     | elapsed: 9.5s
[Parallel(n_jobs=-1)]: Done 434 tasks     | elapsed: 22.5s
[Parallel(n_jobs=-1)]: Done 810 out of 810 | elapsed: 38.9s finished
```

Fitting 5 folds for each of 162 candidates, totalling 810 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
```

```
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed: 2.7s
[Parallel(n_jobs=-1)]: Done 184 tasks     | elapsed: 10.6s
[Parallel(n_jobs=-1)]: Done 434 tasks     | elapsed: 22.3s
[Parallel(n_jobs=-1)]: Done 784 tasks     | elapsed: 45.0s
[Parallel(n_jobs=-1)]: Done 810 out of 810 | elapsed: 45.7s finished
```

Fitting 5 folds for each of 162 candidates, totalling 810 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
```

```
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed: 2.4s
[Parallel(n_jobs=-1)]: Done 184 tasks     | elapsed: 9.1s
[Parallel(n_jobs=-1)]: Done 434 tasks     | elapsed: 22.9s
[Parallel(n_jobs=-1)]: Done 784 tasks     | elapsed: 43.6s
[Parallel(n_jobs=-1)]: Done 810 out of 810 | elapsed: 44.2s finished
```

Mean Recall Score: 0.7104807212391298

Best parameters: {'criterion': 'gini', 'max\_depth': 10, 'max\_features': 'log2', 'min\_samples\_leaf': 5, 'min\_samples\_split': 15}

Best score: 0.7490241520767522

```
In [146]: vif_undersample_dt_grid_search_best_params = vif_undersample_dt_grid_search.best_params_
```

## VIF SMOTE Decision Tree Tuning

```
In [147]: tuned_vif_smote_dt_model_df = wells_df.copy()
```

```
X_train, X_test, y_train, y_test = split_data(tuned_vif_smote_dt_model_df, 'status_group')
```

```
In [148]: tuned_vif_smote_dt_model = DecisionTreeClassifier(random_state=42)
```

```
In [149]: # Create an instance of CustomCrossValidator
tuned_vif_smote_dt_cv = CustomCrossValidator(
    model=tuned_vif_smote_dt_model,
    X=X_train,
    y=y_train,
    sampling_method='smote'
)
```

```
In [150]: vif_smote_dt_grid_search = tuned_vif_smote_dt_cv.run_grid_search(
    param_grid=decision_tree_param_grid,
    features=vif_selected_features,
    cv=StratifiedKFold(n_splits=5)
)
```

Fitting 5 folds for each of 162 candidates, totalling 810 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
```

```
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed: 3.4s
[Parallel(n_jobs=-1)]: Done 184 tasks     | elapsed: 13.8s
[Parallel(n_jobs=-1)]: Done 434 tasks     | elapsed: 34.5s
[Parallel(n_jobs=-1)]: Done 784 tasks     | elapsed: 1.1min
[Parallel(n_jobs=-1)]: Done 810 out of 810 | elapsed: 1.1min finished
```

Fitting 5 folds for each of 162 candidates, totalling 810 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
```

```
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed: 3.3s
[Parallel(n_jobs=-1)]: Done 184 tasks     | elapsed: 12.9s
[Parallel(n_jobs=-1)]: Done 434 tasks     | elapsed: 30.8s
[Parallel(n_jobs=-1)]: Done 784 tasks     | elapsed: 1.0min
[Parallel(n_jobs=-1)]: Done 810 out of 810 | elapsed: 1.0min finished
```

Fitting 5 folds for each of 162 candidates, totalling 810 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
```

```
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed: 4.7s
[Parallel(n_jobs=-1)]: Done 184 tasks     | elapsed: 15.8s
[Parallel(n_jobs=-1)]: Done 434 tasks     | elapsed: 33.8s
[Parallel(n_jobs=-1)]: Done 784 tasks     | elapsed: 56.7s
[Parallel(n_jobs=-1)]: Done 810 out of 810 | elapsed: 58.1s finished
```

Fitting 5 folds for each of 162 candidates, totalling 810 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
```

```
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed: 3.7s
[Parallel(n_jobs=-1)]: Done 184 tasks     | elapsed: 13.2s
[Parallel(n_jobs=-1)]: Done 434 tasks     | elapsed: 35.3s
[Parallel(n_jobs=-1)]: Done 784 tasks     | elapsed: 1.1min
[Parallel(n_jobs=-1)]: Done 810 out of 810 | elapsed: 1.1min finished
```

Fitting 5 folds for each of 162 candidates, totalling 810 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
```

```
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed: 3.6s
[Parallel(n_jobs=-1)]: Done 184 tasks     | elapsed: 14.6s
[Parallel(n_jobs=-1)]: Done 434 tasks     | elapsed: 31.4s
[Parallel(n_jobs=-1)]: Done 784 tasks     | elapsed: 52.4s
[Parallel(n_jobs=-1)]: Done 810 out of 810 | elapsed: 53.2s finished
```

Mean Recall Score: 0.7214271280109246

Best parameters: {'criterion': 'entropy', 'max\_depth': 10, 'max\_features': 'log2', 'min\_samples\_leaf': 4, 'min\_samples\_split': 10}

Best score: 0.7406721731700371

```
In [151]: vif_smote_dt_grid_search_best_params = vif_smote_dt_grid_search.best_params_
```

Oversampling yielded a better scoring model without hyperparameter tuning and had the highest mean recall score amongst all validation folds, the undersampling approach produced the best individual model on the validation set.

Given our business problem, we want to choose the model that performs consistently over a one-time better test.

As such, we have chosen the oversampling model as our best model.

## Fitting the Tuned Best Model

### Fit and Feature Importance

```
In [173]: final_model_df = wells_df.copy()
```

```
X_train, X_test, y_train, y_test = split_data(final_model_df, 'status_group')
```

```
In [174]: final_model = DecisionTreeClassifier(
            random_state=42,
            **vif_oversample_dt_grid_search_best_params
        )
```

```
In [175]: X_train_numeric_scaled_df, X_test_numeric_scaled_df = scale_numeric_features(X_train, X_test)
```

```
In [176]: encoded_train_df, encoded_test_df = encode_categorical_features(X_train, X_test)
```

```
In [177]: x_train_preprocessed_ = combine_features(X_train_numeric_scaled_df, encoded_train_df)
x_test_preprocessed_ = combine_features(X_test_numeric_scaled_df, encoded_test_df)
```

```
In [178]: # Fit the model on the preprocessed training data
final_model.fit(X_train_preprocessed_, y_train)
```

```
Out[178]: DecisionTreeClassifier(max_depth=5, max_features='log2', min_samples_leaf=4,
                                min_samples_split=5, random_state=42)
```

```
In [179]: # Get feature importances
importances = final_model.feature_importances_
```

```
In [180]: # Get column names after preprocessing
preprocessed_columns = X_train_preprocessed_.columns
```

```
In [181]: # Create a DataFrame to hold feature importances with corresponding names
feature_importances = pd.DataFrame({'feature': preprocessed_columns, 'importance': importances})
```


```
In [182]: # Sort features by importance in descending order
feature_importances = feature_importances.sort_values(by='importance', ascending=False)
```

```
In [183]: # Print or visualize the top features
print("Top 10 Most Important Features:")
print(feature_importances.head(10))
```

Top 10 Most Important Features:

	feature	importance
0	amount_tsh	0.449971
259	source_spring	0.097126
240	water_quality_unknown	0.092763
189	extraction_type_unknown	0.066650
59	region_kigoma	0.057295
248	quantity_enough	0.037538
15	funder_other	0.033881
99	region_code_90	0.024713
57	region_iringa	0.024641
231	payment_pay when scheme fails	0.024184

```
In [184.. plt.figure(figsize=(10, 6))
plt.barh(feature_importances['feature'][:10], feature_importances['importance'][:10], align='center')
plt.xlabel('Importance')
plt.title('Top 10 Most Important Features')
plt.gca().invert_yaxis() # Invert y-axis to have the most important feature at the top
plt.show()
```



## Predicting Using the Hold Out Test

```
In [185.. # Predict on the preprocessed test data
y_pred = final_model.predict(X_test_preprocessed)
```


```
In [186.. final_model_recall_score = recall_score(y_test, y_pred)
print(f"Recall on test set: {final_model_recall_score:.4f}")
```

Recall on test set: 0.7476

```
In [187.. # Compute confusion matrix
cm = confusion_matrix(y_test, y_pred)
```

```
In [188.. # Plot confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False)
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()
```





Confusion Matrix Breakdown:

- True Positive (TP): Correctly predicting a well as "Does not need attention" (0).
- False Positive (FP): Incorrectly predicting a well as "Needs attention" (1) when it actually does not need attention.
- True Negative (TN): Correctly predicting a well as "Needs attention" (1).
- False Negative (FN): Incorrectly predicting a well as "Does not need attention" (0) when it actually needs attention.

There is a slight improvement in recall score on the test data compared to the training data when using the oversampling model with VIF-selected features.

This suggests that our model has generalized well to unseen data. Which given the drop in recall score from the train to the hyperparameter tuned train alleviates some of the fears around overfitting.

The recall score on the training data was 0.7414, while the test data scored 0.7476, indicating that the model can effectively identify wells needing attention.

## Recommendations

### Business Recommendations

#### Prioritize wells in the northwest and southeast of the country

Out of the top twenty regions that have wells in need of attention, these regions are in the top 20.

If you could only prioritize one region, start in the northwest of the country.

Northwest:

- Kagera
- Kigoma
- Shinyanga
- Mwanza
- Mara

Southeast:

- Lindi
- Mtwara
- Iringa

#### For new wells, look to improve or use the latest version of gravity extraction pumps

Gravity extraction pumps are the most prevalent extraction type used in wells that need attention.

This is counterintuitive as gravity powered well pump should not be as susceptible to wear and tear as a motorized or centrifugal pumps.

It appears as though the construction quality and materials used could be playing a role here.

# Predictive Recommendation

All data provided was surveyed by GeoData Consultants Ltd. Our prediction model can be used in place of consultants, cutting costs.

## Benefit of the model

- Can assist in efficient resource allocation and proactive maintenance planning.
  - **Note**
    - The model is based on historic results, will not be apt for use during natural disasters but can be used before-hand to identify high priority regions.

In [ ]: